

Towards optimal synchronous counting

Christoph Lenzen

MPI for Informatics

Joel Rybicki

MPI for Informatics
Aalto University

Jukka Suomela

Aalto University

PODC 2015

July 23

Focus on fault-tolerance

Fault-tolerant ***co-ordination*** primitives:

- **permanent** failures (*Byzantine faults*)
- **transient** failures (*self-stabilisation*)

Focus on fault-tolerance

Fault-tolerant ***co-ordination*** primitives:

- **permanent** failures (*Byzantine faults*)
- **transient** failures (*self-stabilisation*)

Find solutions that are

- **fast** to recover
- **space** and **communication**-efficient

Our contribution

A **deterministic** *round counter* with


- **high** resilience
- **optimal** recovery time
- **low** space/message complexity

Model of computing

n state machines

Synchronous rounds:

1. broadcast
2. receive
3. update state



Model of computing


n state machines

Synchronous rounds:

1. broadcast

2. receive

3. update state



Model of computing


n state machines

Synchronous rounds:

1. broadcast

2. receive

3. update state



Model of computing


n state machines

Synchronous rounds:

1. broadcast

2. receive

3. update state



Algorithm **A** maps a *vector* of states
to a new state!

Complexity measures

Time complexity: #rounds

\approx “recovery time”

Space complexity: \log #states

\approx complexity of the circuit

\approx number of bits broadcast per node

On failures



our adversary

Transient failures


n state machines

arbitrary initial states



chosen by adversary!


= self-stabilisation



Byzantine failures

n state machines


f Byzantine failures



Byzantine failures

n state machines


f Byzantine failures



Byzantine failures

n state machines


f Byzantine failures



Byzantine failures

n state machines


f Byzantine failures



Byzantine failures

n state machines

f Byzantine failures



**Non-faulty nodes can observe
different states for the system!**

Counting mod c

3-counting

0	1	2	0	1	2
---	---	---	---	---	---

increment counter $+1 \bmod c$

Synchronous counting

Counting

1

0 1 2 0 1 2

②


0 1 2 0 1 2

③

0 1 2 0 1 2

④


0 1 2 0 1 2



Global pulse


Synchronous counting

Stabilisation Counting



Synchronous counting

Stabilisation Counting



A related problem:

Consensus

Input: private bit

Output: agreement

0

1



1

A related problem:

Consensus

Input: private bit


Output: agreement

0

1



1



1

1



1

A related problem:

Consensus

Input: private bit


Output: agreement

0

1



1



0

0



0

A related problem:

Consensus

Input: private bit


Output: agreement

0

0



0



0

0




0

A related problem:

Consensus

Input: private bit

Output: agreement



Reduction from consensus

Given a 2-counting algorithm
A that stabilises in ***t*** rounds



consensus solvable in ***t*** rounds

Dolev et al. (2013)

Consensus lower bounds*

Resilience

$$f < n/3$$

Pease et al. (1980)

Time

At least f rounds to reach agreement

Fischer & Lynch (1982)

***deterministic**

Prior work on 2-counting

Resilience

Time

State bits

D. Dolev &
Hoch 2007 *

$$f < n/3$$

$$O(f)$$

$$O(f \log f)$$

***deterministic**

Prior work on 2-counting

Resilience

Time

State bits

D. Dolev &
Hoch 2007 *

$$f < n/3$$

$$O(f)$$

$$O(f \log f)$$

S. Dolev &
Welch 2004

$$f < n/3$$

$$2^{2(n-f)}$$

$$O(1)$$

***deterministic**

Prior work on 2-counting

Resilience

Time

State bits

D. Dolev &
Hoch 2007 *

$$f < n/3$$

$$O(f)$$

$$O(f \log f)$$

S. Dolev &
Welch 2004

$$f < n/3$$

$$2^{2(n-f)}$$

$$O(1)$$

Ben-Or *et al.*
2008

$$f < n/3$$

$$O(1)$$

$$n^{O(1)}$$

***deterministic**

Current state

Resilience

Time

State bits

D. Dolev &
Hoch 2007 *

$$f < n/3$$

$$O(f)$$

$$O(f \log f)$$

S. Dolev &
Welch 2004

$$f < n/3$$

$$2^{2(n-f)}$$

$$O(1)$$

Ben-Or *et al.*
2008

$$f < n/3$$

$$O(1)$$

$$n^{O(1)}$$

Our work*

$$f = n^{1-o(1)}$$

$$O(f)$$

$$o(\log^2 f)$$


***deterministic**

Counting vs consensus



Counting

?




Consensus

Counting vs consensus




Counting




Consensus

Solve consensus to agree on counters

Counting vs consensus



Counting




Consensus

Use a *c*-counter as a round counter;
execute a consensus algorithm

Counting vs consensus



?




Counting

Consensus

Counting vs consensus



!




Counting

Consensus


Solution: counters that
work *once in a while*

Counting once in a while

①	2	0	0	1	2	2	0	0	0	1	2
②	2	1	0	1	2	0	1	2	0	1	2
	?	?	?	?	?	?	?	?	?	?	?
④	0	1	0	1	2	0	0	1	0	1	2
	Arbitrary Counting					Arbitrary Counting					


Counting once in a while

Use *proper* counters with **low resilience**, to count *once in a while* with **high resilience**!



Counting

low resilience



Counting
once in a while
high resilience

Clock for consensus

If we can count *once in a while* with
high resilience..



then we can execute a **highly-resilient
consensus** algorithm (*phase king*)


Berman et al. (1989)

Agree on a new counter

Use consensus protocol
with **high resilience**, to agree on a
new proper counter!




High-level idea




Counting
low resilience

High-level idea




Counting
low resilience




Counting
once in a while
high resilience


High-level idea




Counting
low resilience




Counting
once in a while
high resilience




Consensus
high resilience



High-level idea




Counting
low resilience




Counting
once in a while
high resilience




Consensus
high resilience




Counting
high resilience




Rinse and repeat



**low
resilience**



**high
resilience**




***higher*
resilience**

The boosting theorem

More formally...

The boosting theorem




Algorithm **A**

n nodes



f faults

The boosting theorem



Algorithm **A**

n nodes
 f faults




Algorithm **B**

$$N = kn$$

$$F \approx (f + 1)k/2$$

The boosting theorem




$$F \approx (f + 1)k/2$$

Stabilisation time: $T(\mathbf{B}) = T(\mathbf{A}) + O(Fk^k)$

Space complexity: $S(\mathbf{B}) = S(\mathbf{A}) + O(\log C)$

C = new counter range

Boosting resilience



Boost resilience recursively
while keeping **time** and **space**
complexity *small enough*

Main result

Resilience

$$f = n^{1-o(1)}$$

Stabilisation time

$$O(f)$$

Space complexity

$$O(\log^2 f / \log \log f)$$

Main result

Resilience

$$f = n^{1-o(1)}$$

Stabilisation time

$$O(f)$$

Space complexity

$$O(\log^2 f / \log \log f)$$

Thanks!