

# Self-stabilising Byzantine clock synchronisation is almost as easy as consensus

**Christoph Lenzen**

Max Planck Institute for Informatics

**Joel Rybicki**

University of Helsinki





**DISC 2017**

October 17, 2017

Vienna, Austria







# Clock synchronisation



**Question:**

How to efficiently synchronise a network of  
**imprecise clocks** in a **fault-tolerant** manner?

# Clock synchronisation




## Question:

How to efficiently synchronise a network of **imprecise clocks** in a **fault-tolerant** manner?

Real clocks exhibit **drift**:  
even if started at the same  
time, values eventually differ

Distributed systems are prone  
to both **transient faults**  
and **permanent faults**


# Model of computation



- $n$  nodes with **local clocks**

**Bounded clock drift:**  
fastest clock progresses  
**at most factor  $\vartheta > 1$  faster**  
than the slowest clock

# Model of computation




- $n$  nodes with **local clocks**
- arbitrary initial state

## **Self-stabilisation:**

Initial clock values and local state are arbitrary


# Model of computation



- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes

**Arbitrary misbehaviour:**  
can crash, lose messages,  
send different misinformation  
to correct nodes, etc.


# Model of computation



- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes

**Arbitrary misbehaviour:**  
can crash, lose messages,  
send different misinformation  
to correct nodes, etc.

# Model of computation




- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes

**Arbitrary misbehaviour:**  
can crash, lose messages,  
send different misinformation  
to correct nodes, etc.




# Model of computation



- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes

**Arbitrary misbehaviour:**  
can crash, lose messages,  
send different misinformation  
to correct nodes, etc.


# Model of computation



- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes
- bounded delay communication

**Message delay:**  
transmitting a messages takes  
**between 0 and  $d$  time units**


# Model of computation



- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes
- bounded delay communication

**Message delay:**  
transmitting a messages takes  
**between 0 and  $d$  time units**


# Model of computation



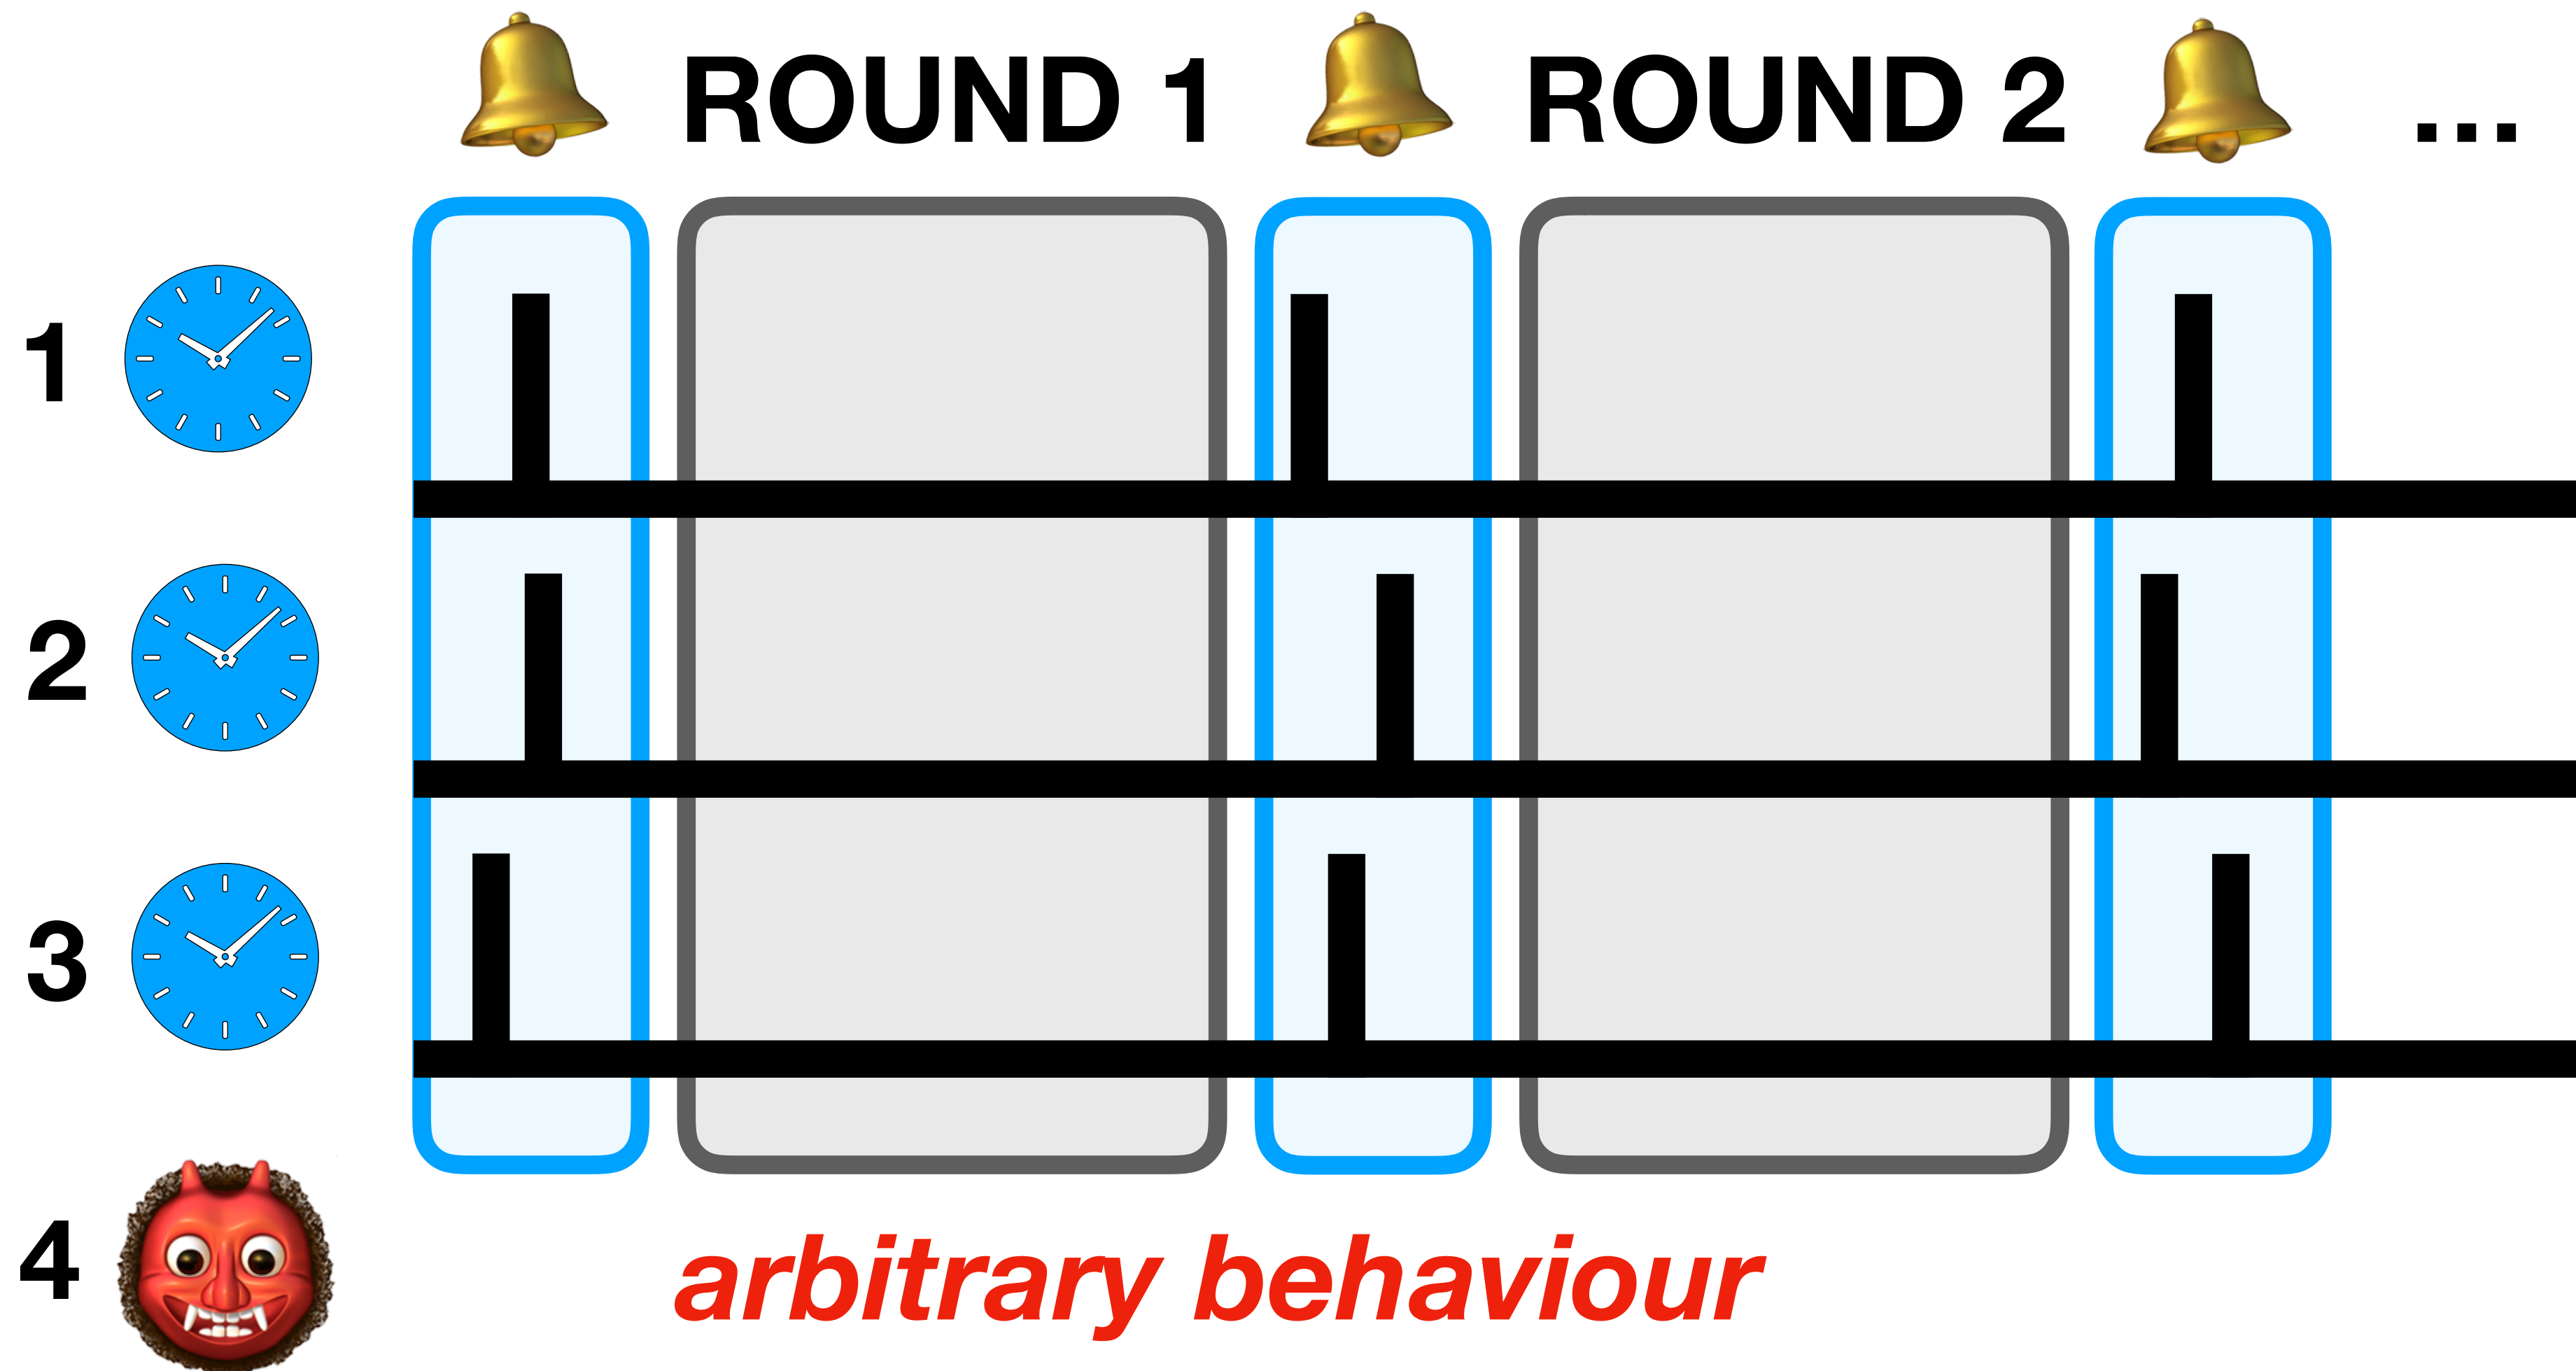
- $n$  nodes with **local clocks**
- arbitrary initial state
- $f < n/3$  **Byzantine faulty** nodes
- bounded delay communication

**Message delay:**  
transmitting a messages takes  
**between 0 and  $d$  time units**

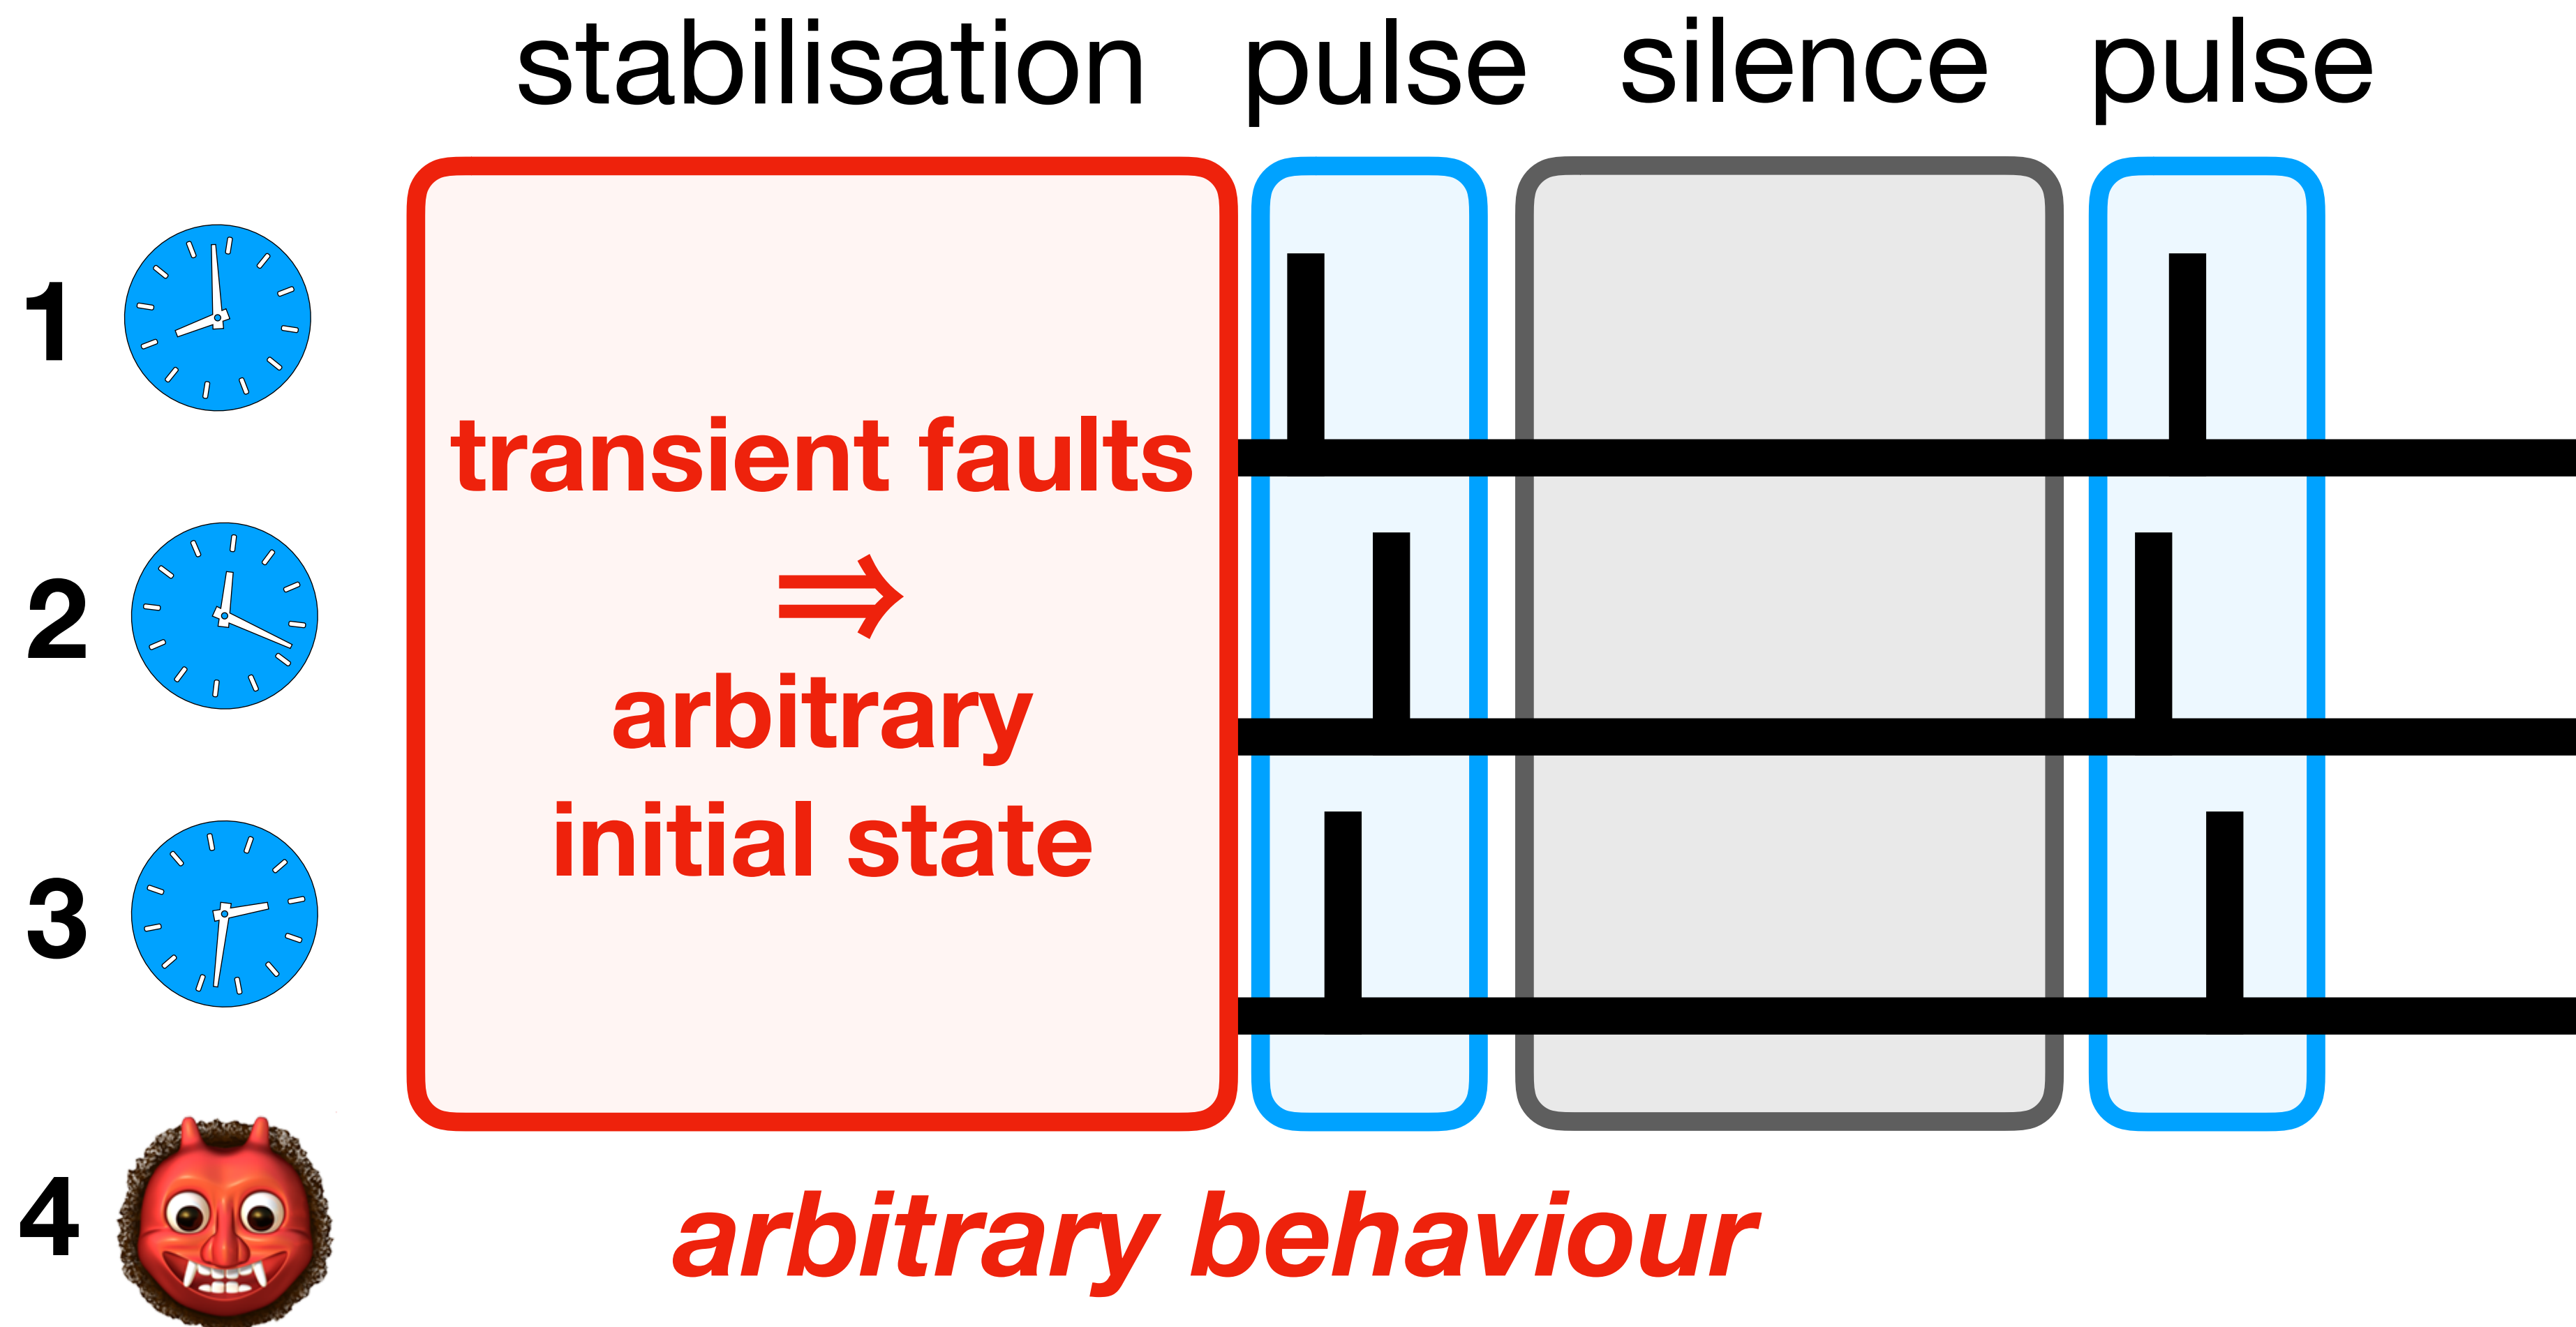
# Pulse synchronisation



# Gives the synchronous model

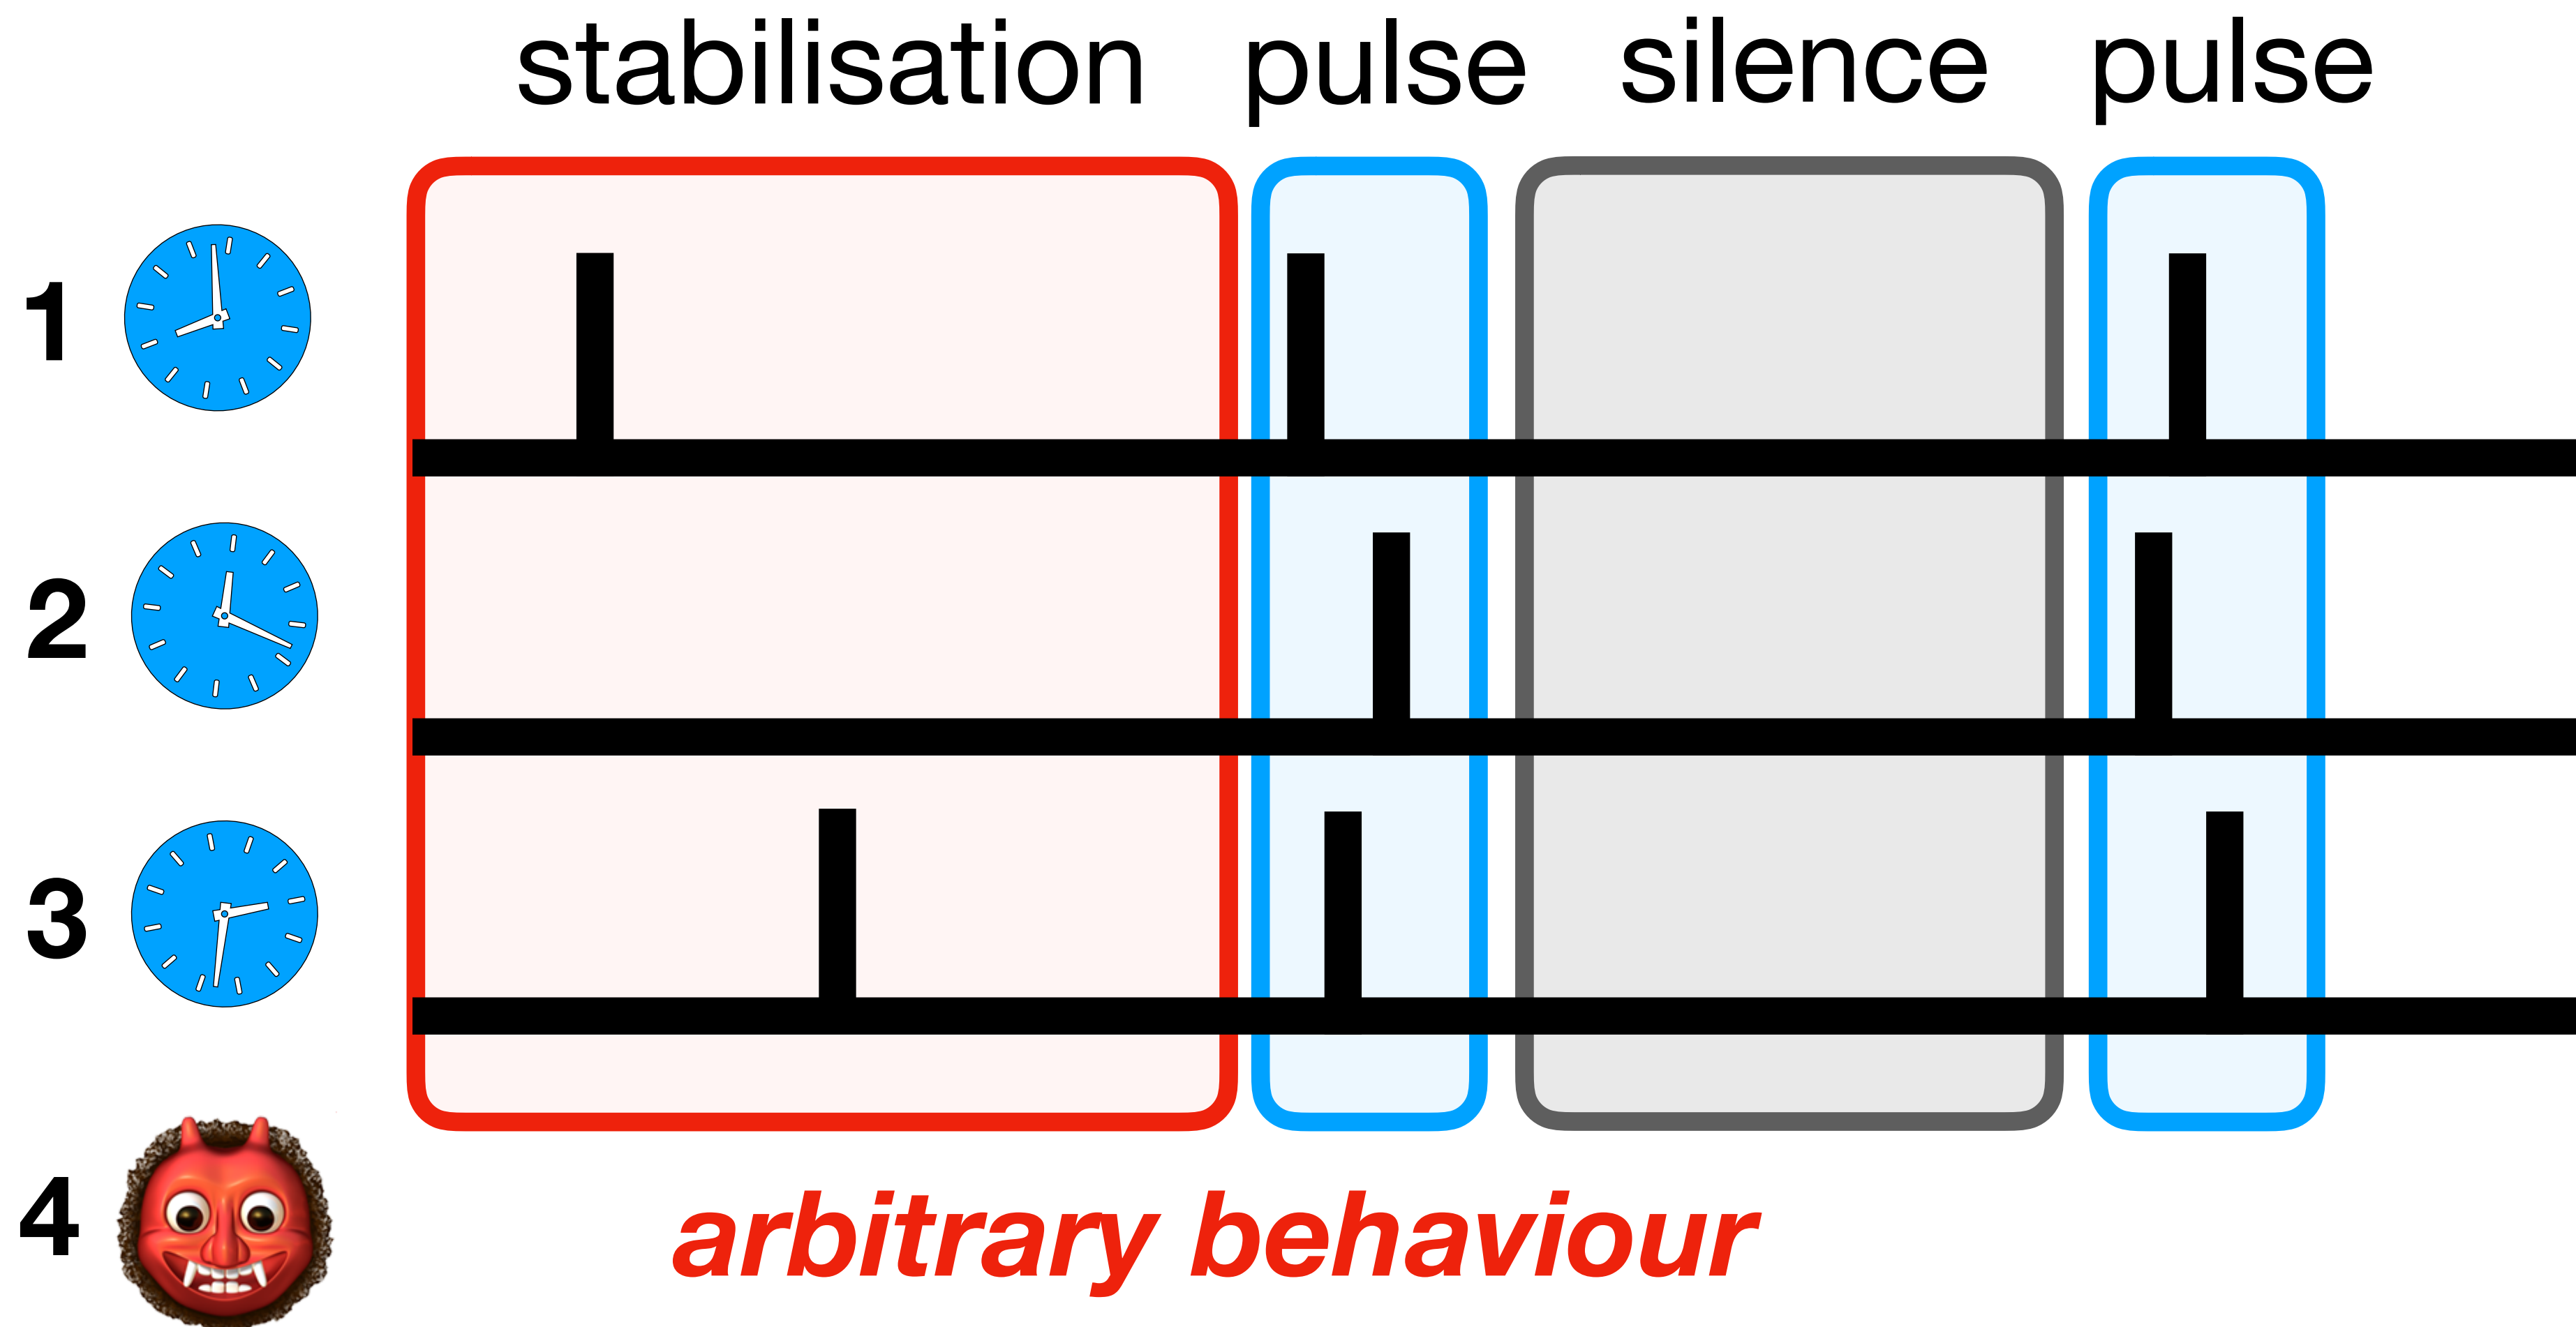


# Self-stabilising pulse synchronisation






# Self-stabilising pulse synchronisation





# Complexity measures

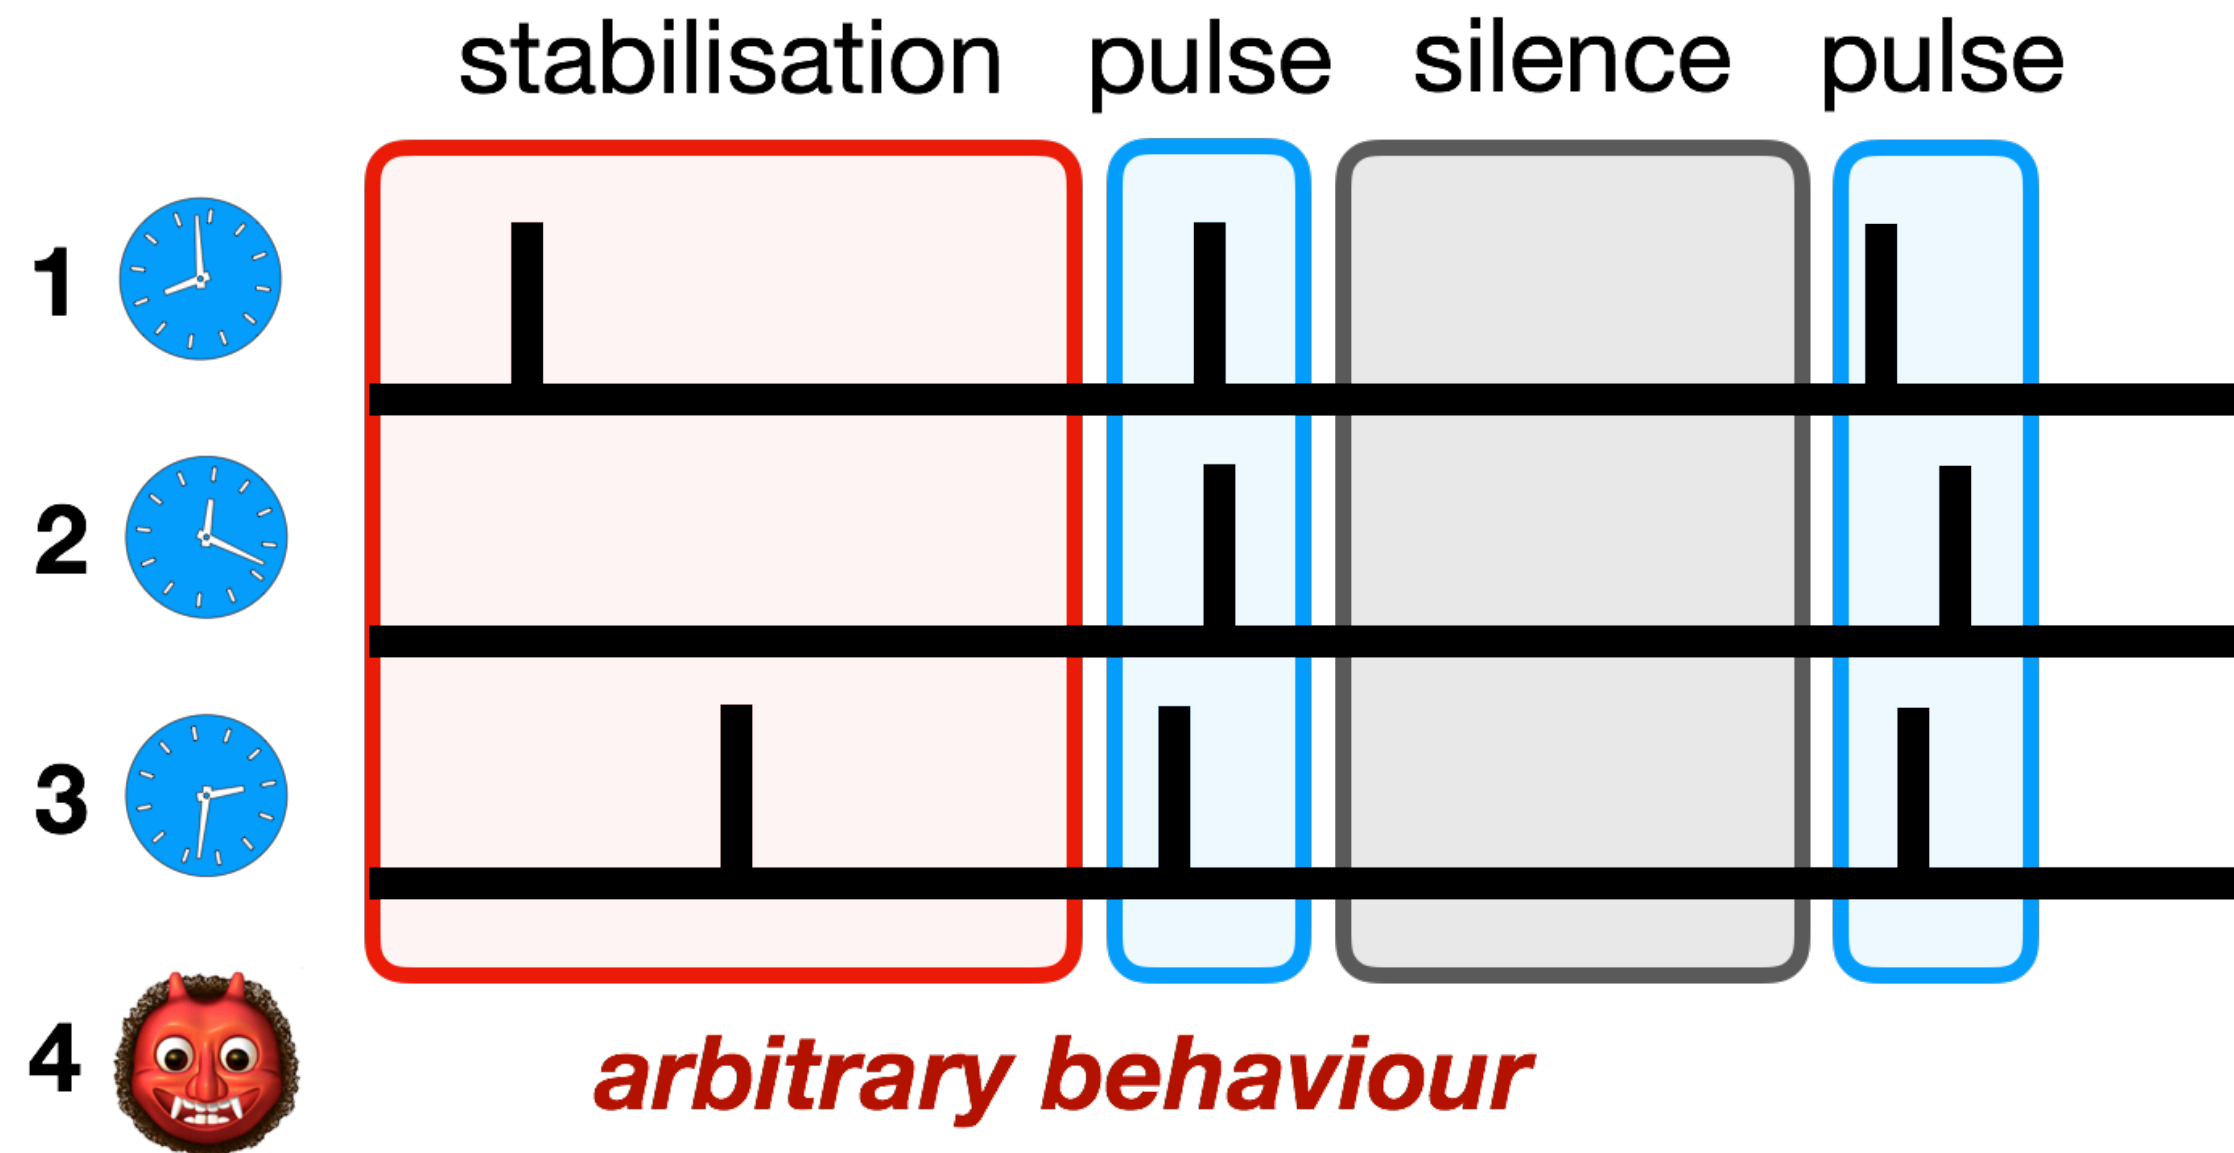


## Stabilisation time:

how many time units it takes before the correct nodes start pulsing synchronously

- $n$  nodes
- $f < n/3$  faults
- drift parameter  $\vartheta = O(1)$
- message delay  $d = O(1)$

# Complexity measures



## Stabilisation time:






how many time units it takes before the correct nodes start pulsing synchronously

## Bandwidth (per node):

how many bits per time unit a node broadcasts


- $n$  nodes
- $f < n/3$  faults
- drift parameter  $\vartheta = O(1)$
- message delay  $d = O(1)$

# Pulse synchronisation: status

authors	appears in		stabilisation time	bandwidth
S. Dolev & Welch	JACM 2004		$\exp(O(f))$	$O(1)$
Dalot, D. Dolev & Parnas	SSS 2003		$O(f^3)$	$O(\log f)$
D. Dolev & Hoch	SSS 2007		$O(f)$	$O(f \log f)$
D. Dolev, Függer, Lenzen & Schmid	JACM 2014		$O(f)$	$O(1)$
this work		 	$\text{polylog } f$ $O(\log f)$ $O(f)$	$\text{polylog } f$ $\text{poly } f$ $O(\log f)$

# Our result

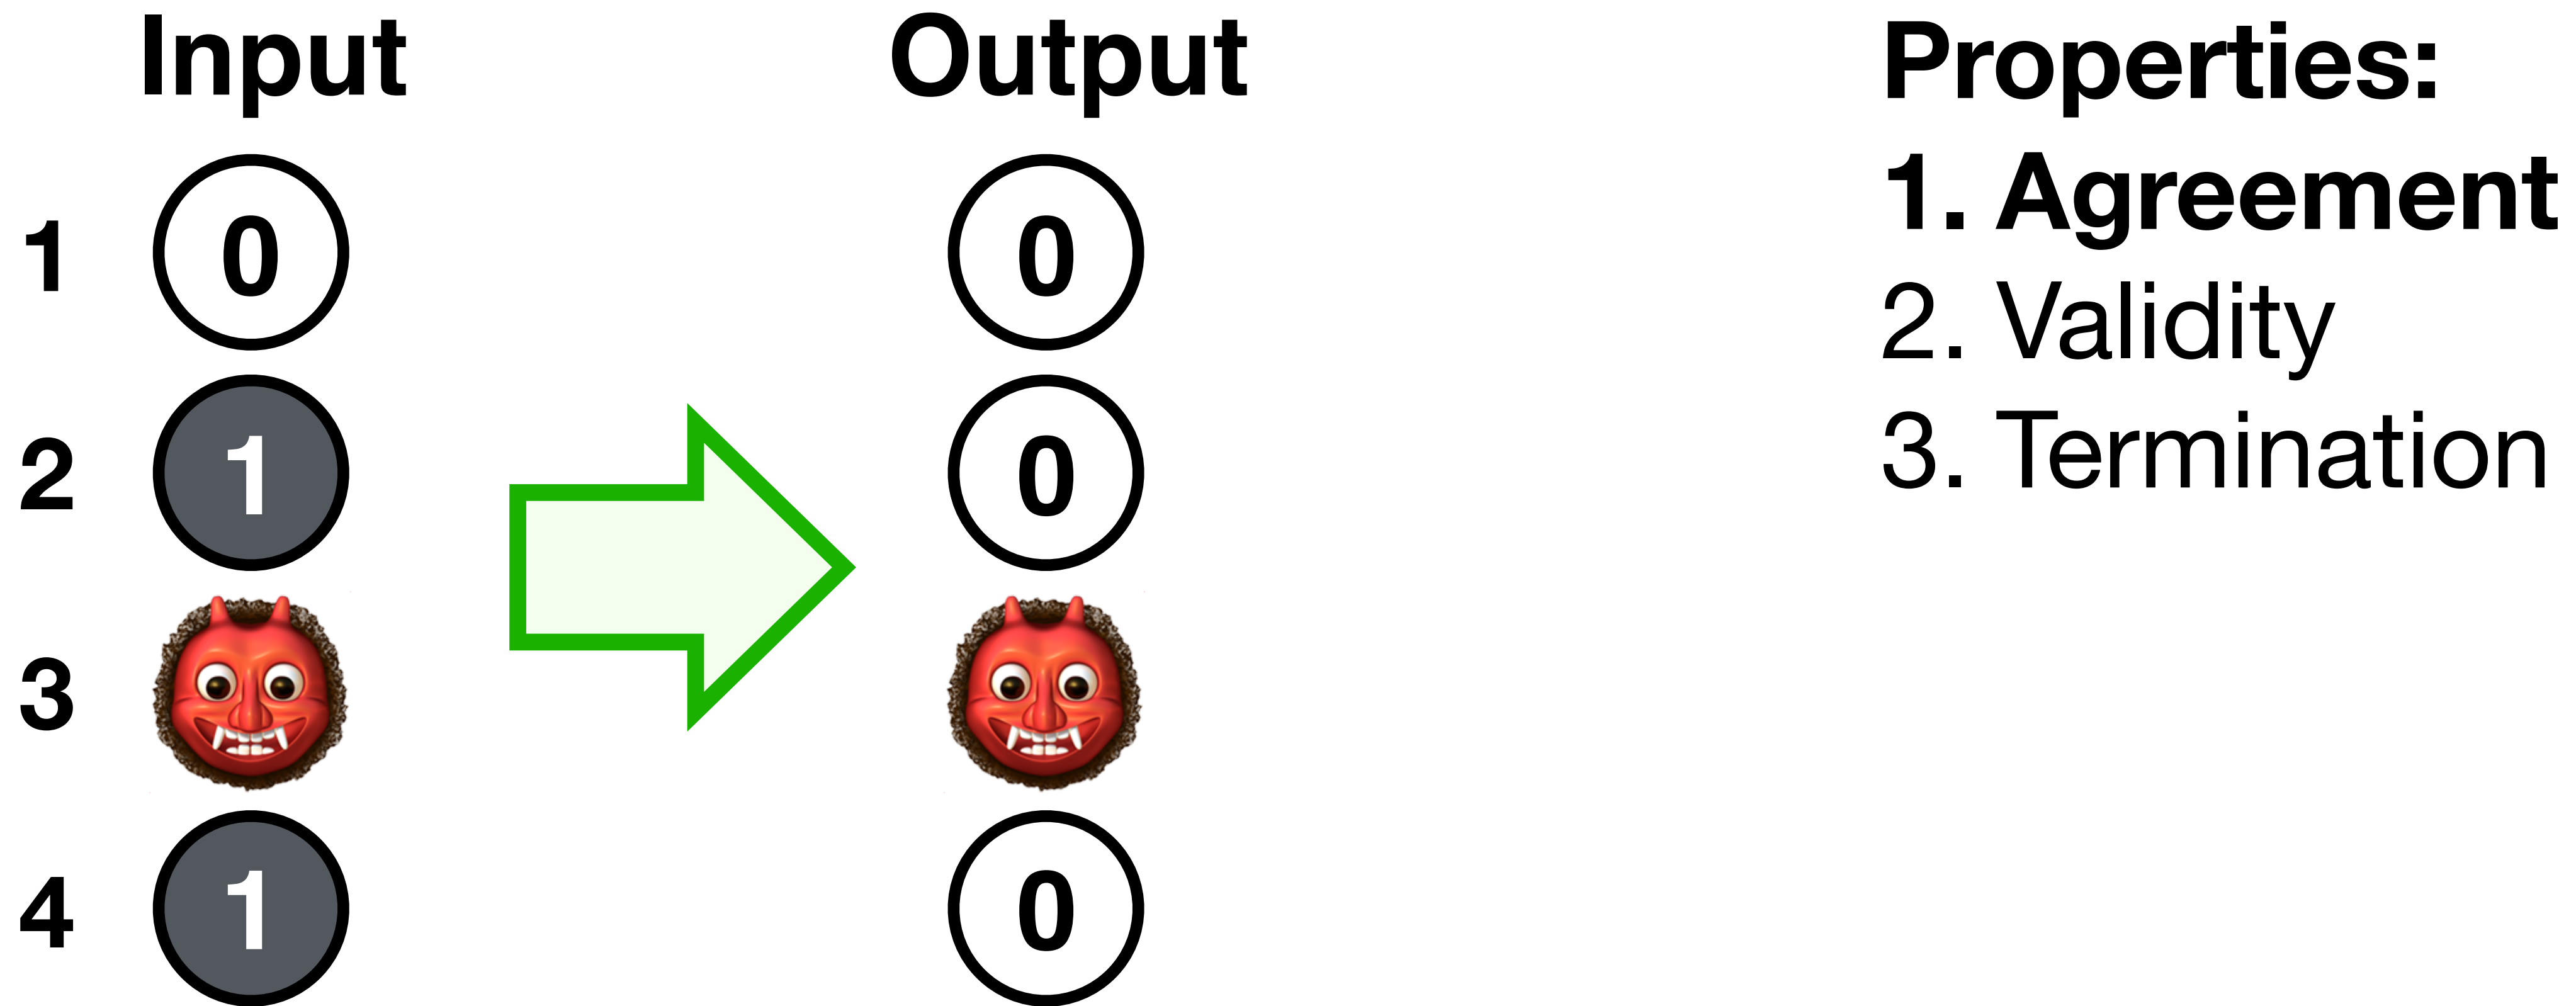
**Synchronous BF consensus**



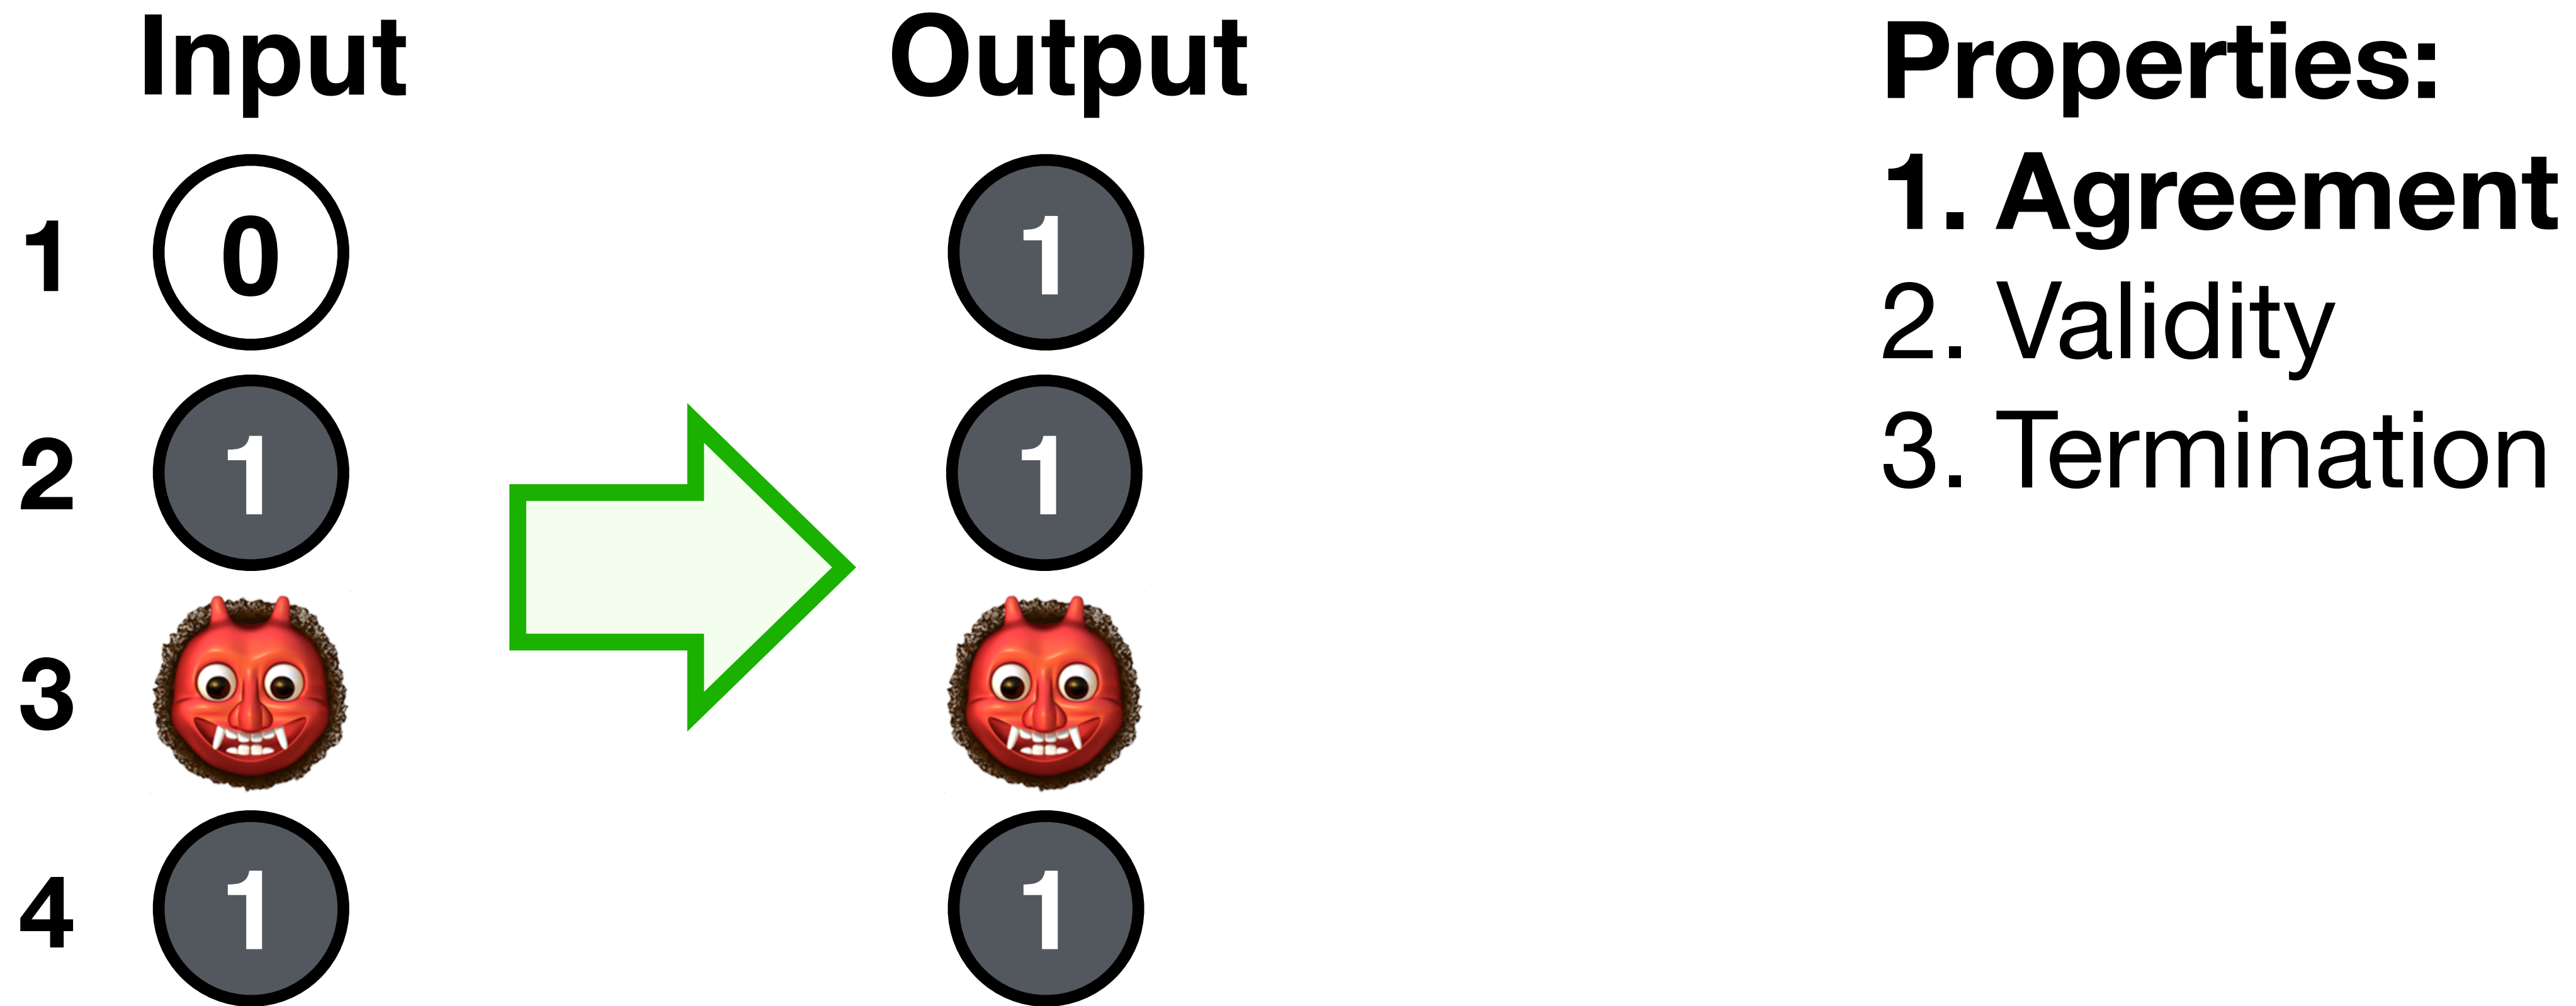
**SS + BF pulse synchronisation  
in the bounded delay model**

(\*) At most **logarithmic** overheads.

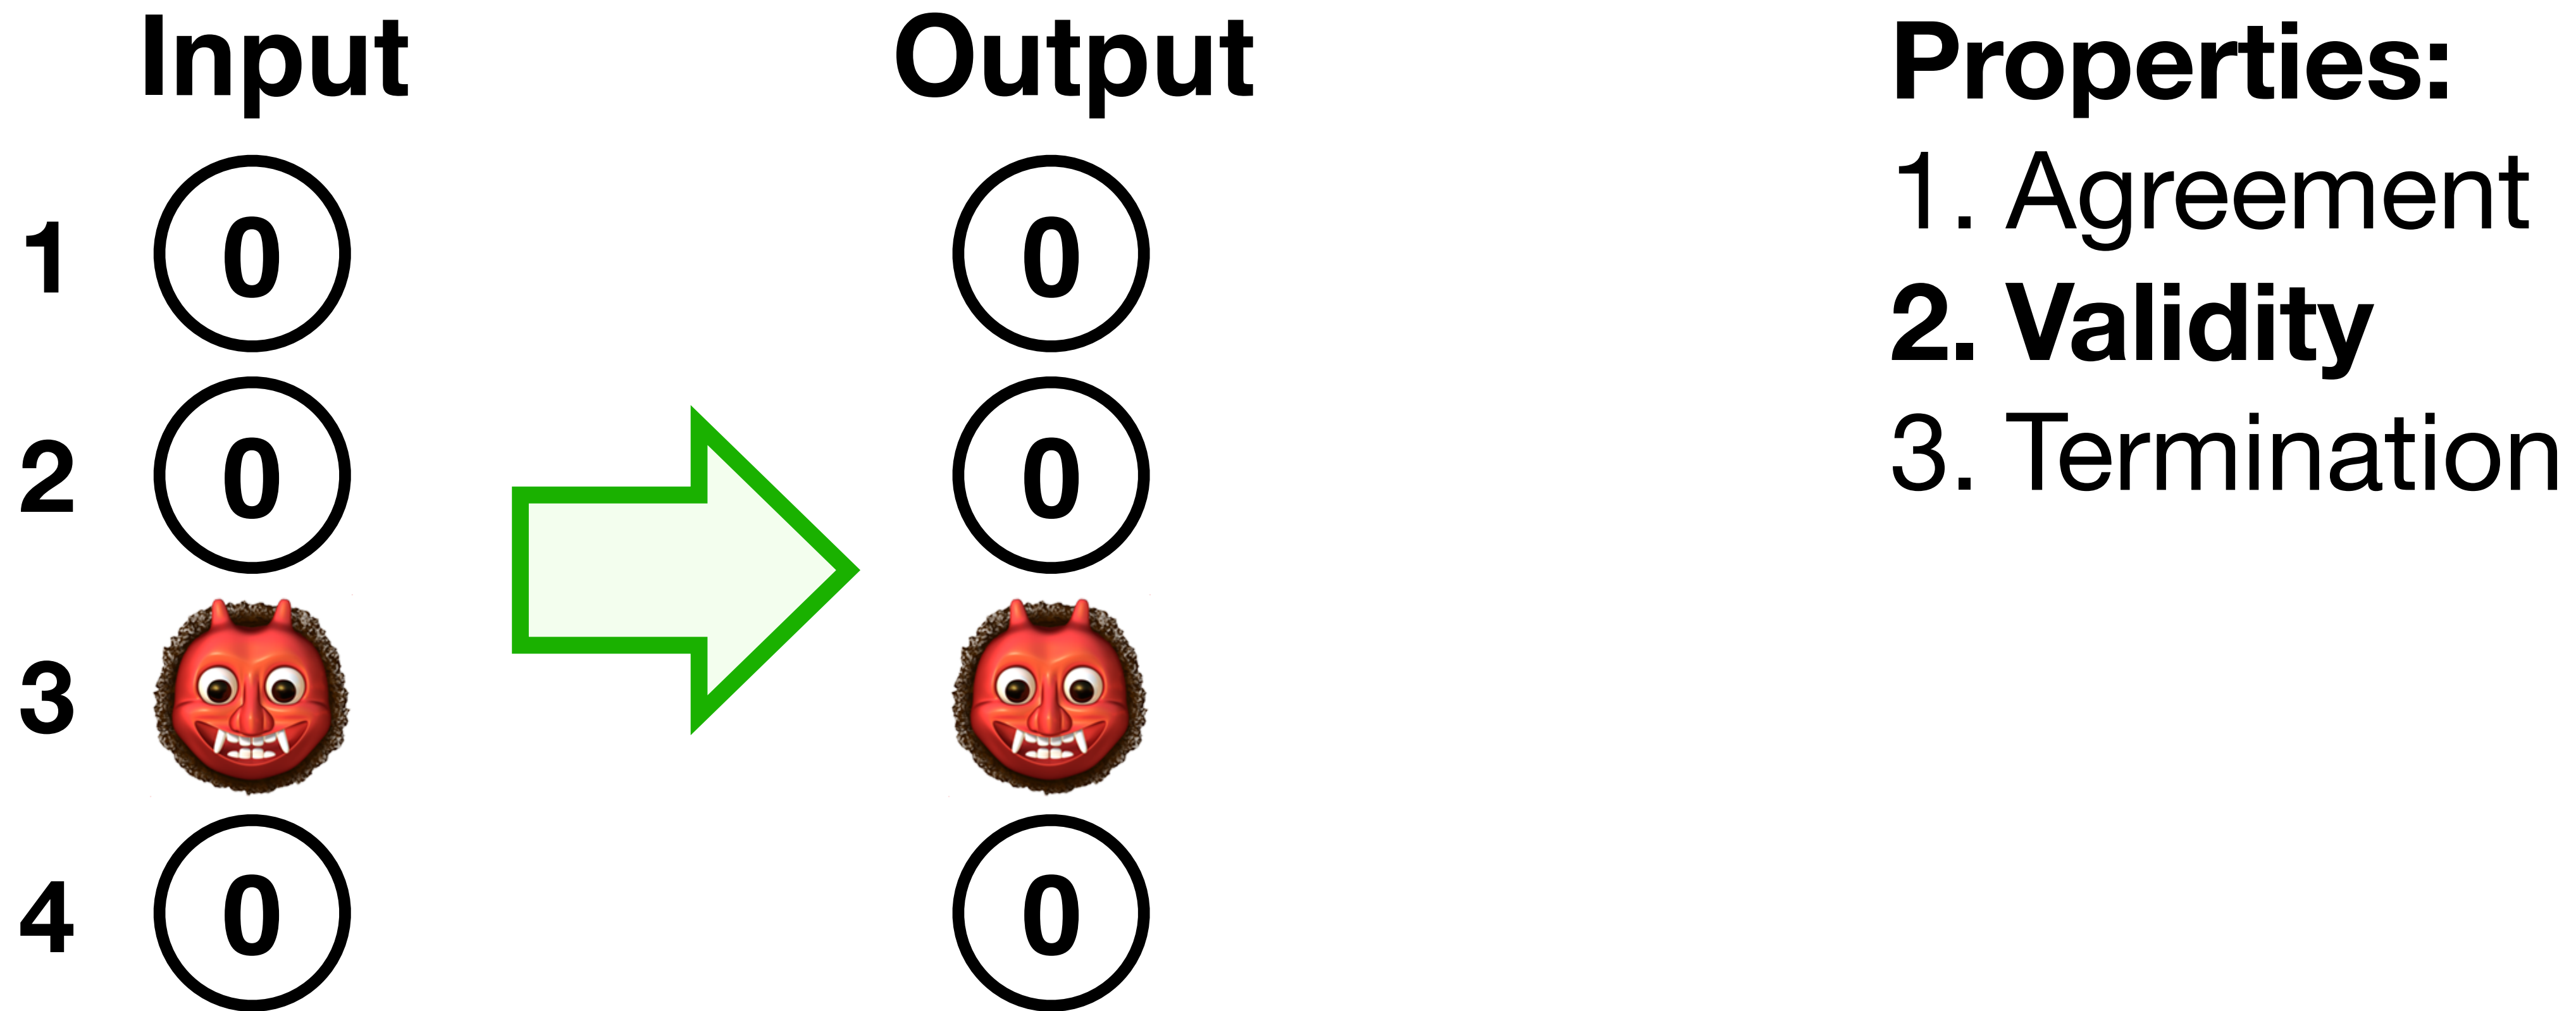
# Synchronous binary consensus



# Synchronous binary consensus

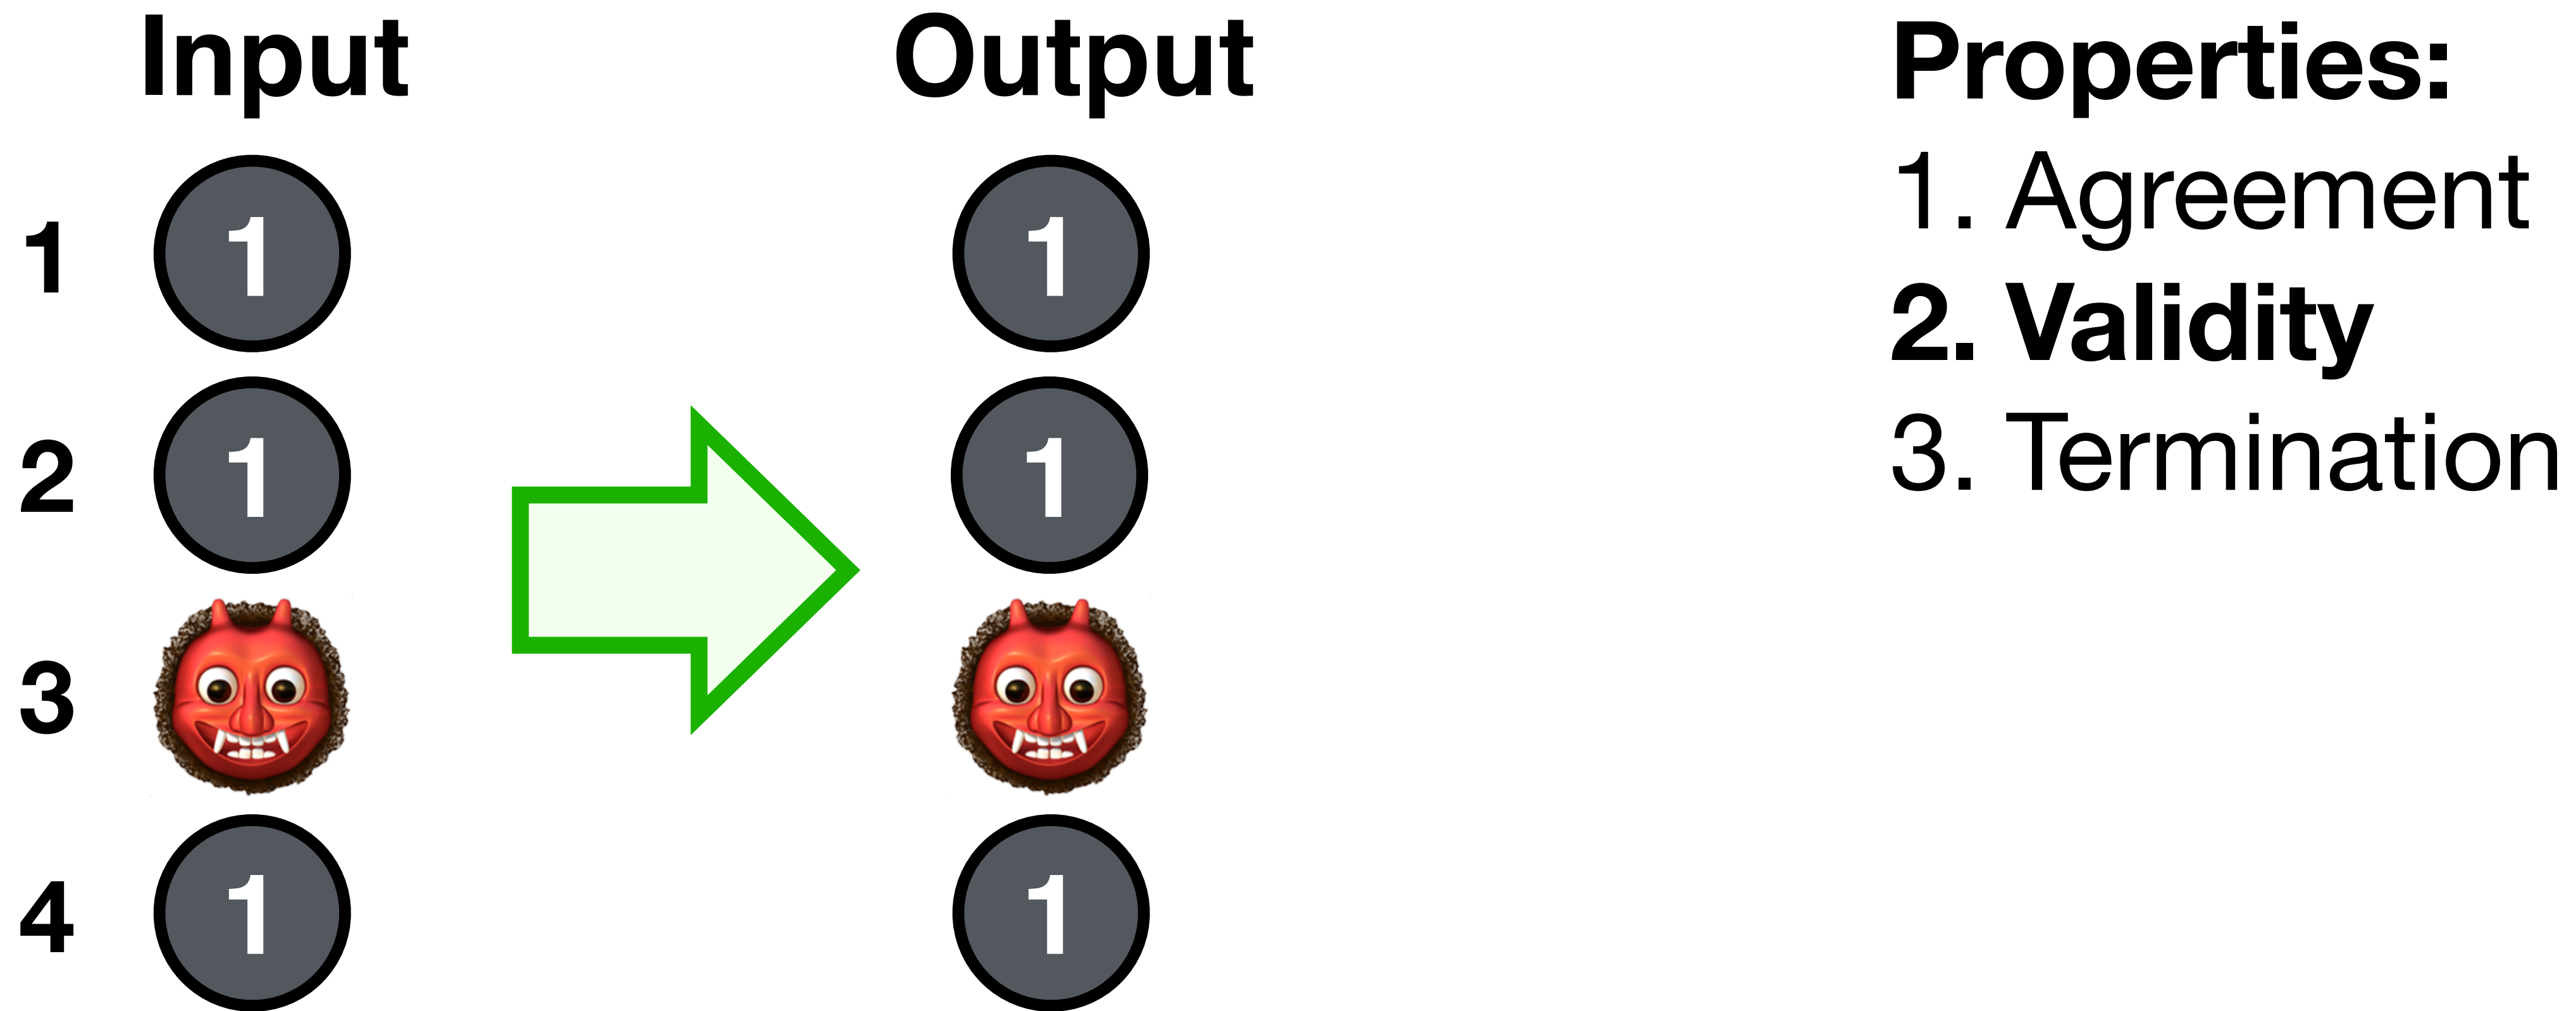


# Synchronous binary consensus






# Synchronous binary consensus





# Synchronous binary consensus



## Properties:


1. Agreement
2. Validity
3. Termination

## Synchronous consensus:

- no self-stabilisation
- **synchronous** communication
- terminates in  $T$  rounds

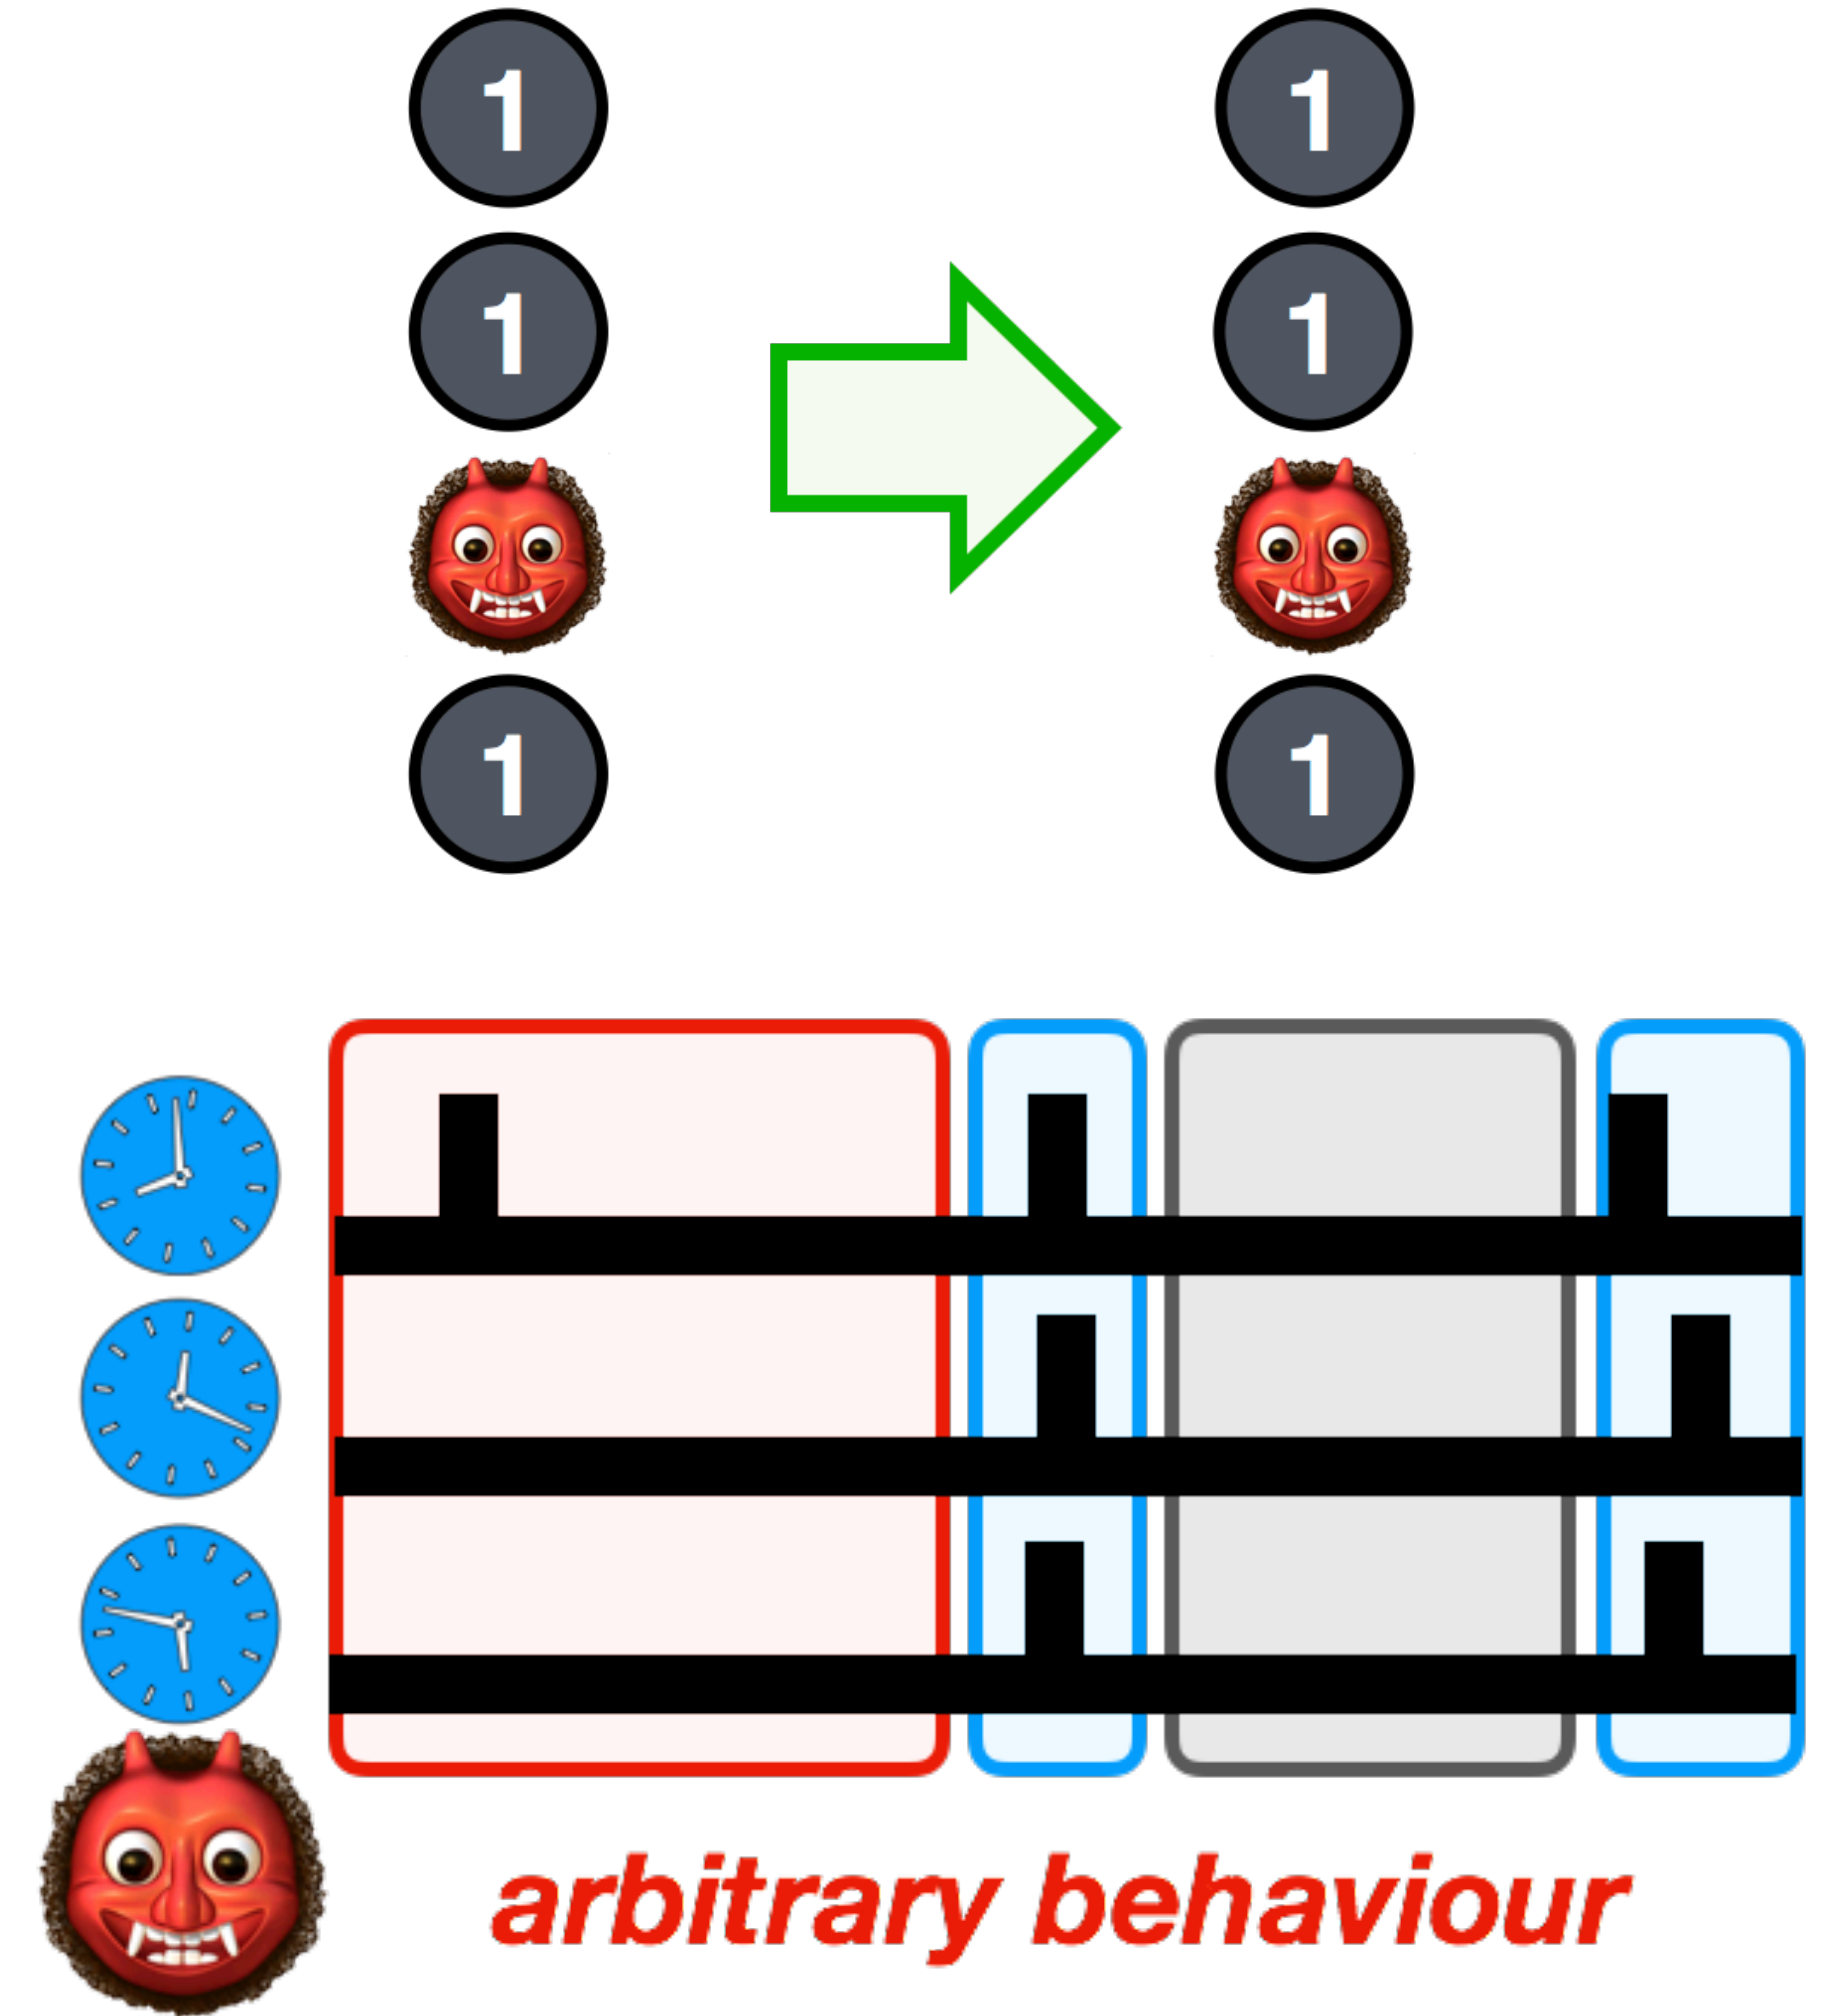
# Our result

**Synchronous BF consensus**



**SS + BF pulse synchronisation  
in the bounded delay model**




(\*) At most **logarithmic** overheads.



# Consensus ➡ pulse sync.

**consensus**

**pulse sync**

source		rounds	bits/round	stabilisation	bandwidth
Berman, Garay & Perry '1992		$O(f)$	$O(1)$	$O(f)$	$O(\log f)$
King & Saia '2011		$\text{polylog}(f)$	$\text{polylog}(f)^*$	$\text{polylog}(f)$	$\text{polylog}(f)$
Feldman & Micali '1988		$O(1)$	$\text{poly}(f)$	$O(\log f)$	$\text{poly}(f)$

# Approach: resilience boosting

Adapt ideas and techniques from prior work on *synchronous counting* (digital clock synchronisation):

- Lenzen & R. (SSS 2016)
- Lenzen, R. & Suomela (SICOMP 2017)


**Difficulty:** translating techniques from the synchronous model to bounded-delay model with clock drift

# Resilience boosting: idea

Given:

Pulser  $A_0$


- $n_0 = 2$  nodes
- $f_0 = 0$  resilience



+

Pulser  $A_1$

- $n_1 = 2$  nodes
- $f_1 = 0$  resilience




# Resilience boosting: idea

Given:

Pulser  $A_0$


- $n_0 = 2$  nodes
- $f_0 = 0$  resilience



+

Pulser  $A_1$


- $n_1 = 2$  nodes
- $f_1 = 0$  resilience



Result:

Pulser  $A$

- $n = n_0 + n_1$  nodes
- $f = f_0 + f_1 + 1$  resilience



# Resilience boosting: idea

Given:

Pulser  $A_0$

- $n_0 = 2$  nodes
- $f_0 = 0$  resilience








+

Pulser  $A_1$

- $n_1 = 2$  nodes
- $f_1 = 0$  resilience



*arbitrary behaviour*







Result:

Pulser  $A$

- $n = n_0 + n_1$  nodes
- $f = f_0 + f_1 + 1$  resilience



*arbitrary behaviour*






# Resilience boosting: idea

Given:

Pulser  $A_0$


- $n_0 = 2$  nodes
- $f_0 = 0$  resilience



+

Pulser  $A_1$


- $n_1 = 2$  nodes
- $f_1 = 0$  resilience



Result:

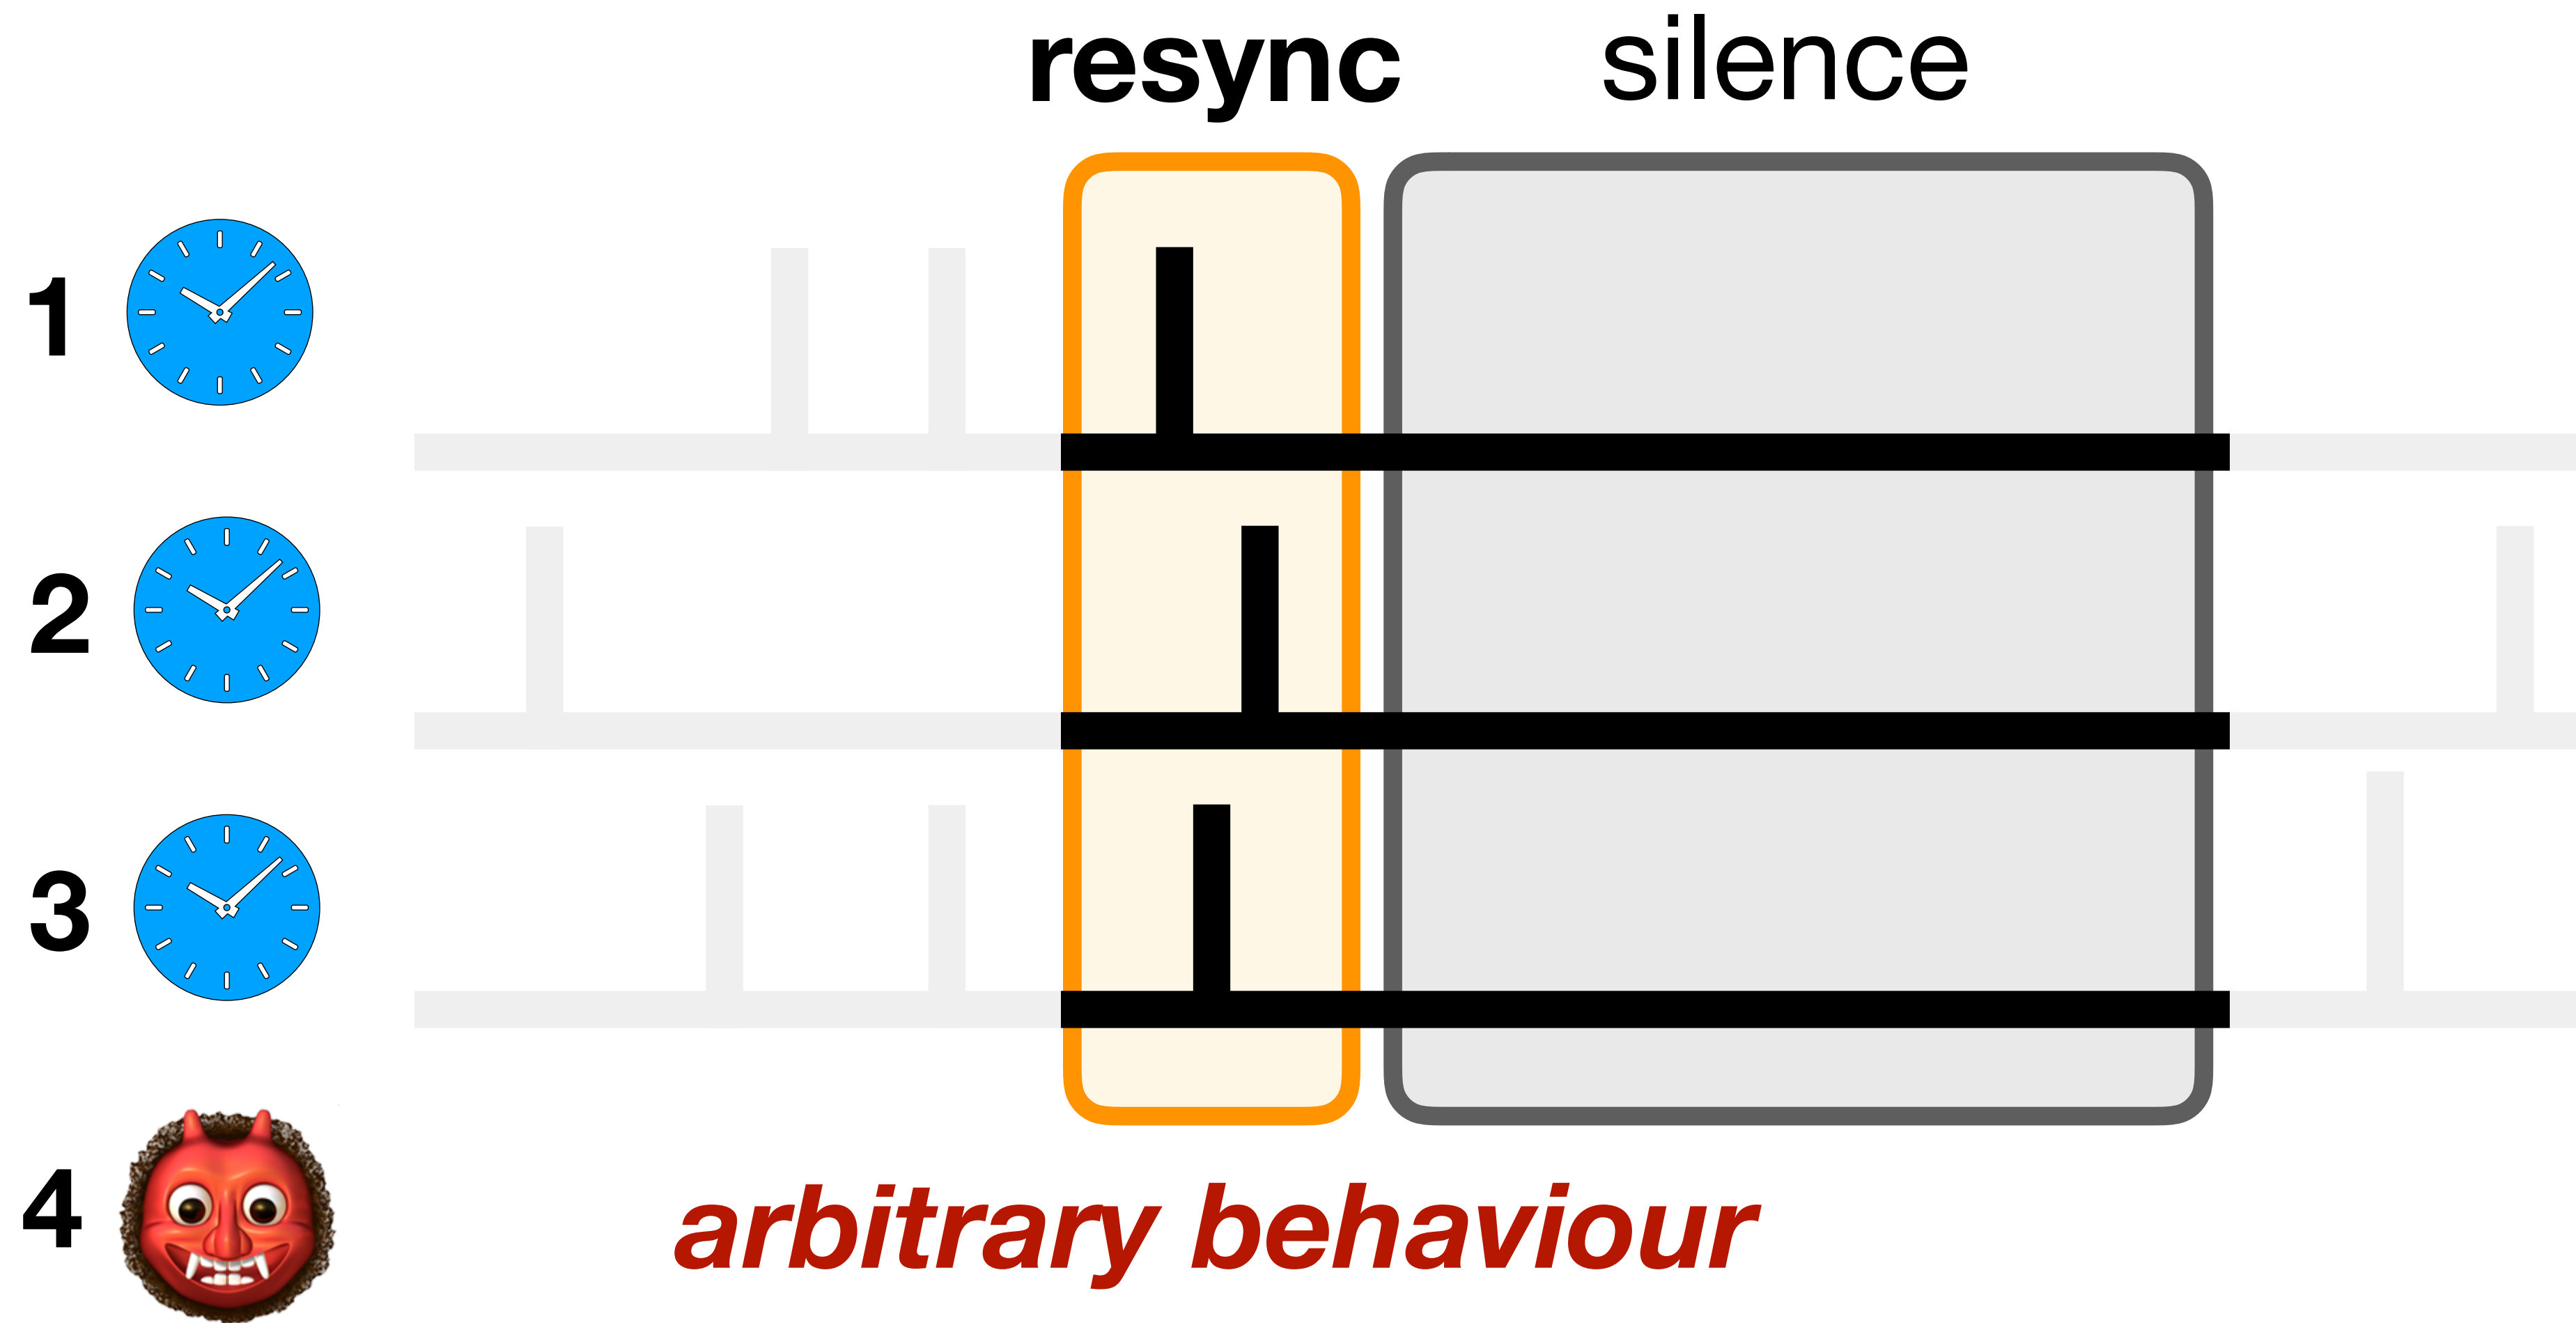
Pulser  $A$

- $n = n_0 + n_1$  nodes
- $f = f_0 + f_1 + 1$  resilience







# Ingredient: resynchronisation pulse



# Ingredient: resynchronisation pulse



# Ingredient: resynchronisation pulse




# Constructing resync. algorithms

Given:

Pulser  $A_0$


- $n_0 = 2$  nodes
- $f_0 = 0$  resilience



+

Pulser  $A_1$

- $n_1 = 2$  nodes
- $f_1 = 0$  resilience




Result:

Resynchroniser **B**

- $n = n_0 + n_1$
- $f = f_0 + f_1 + 1$

Pulser  $A_0$

Pulser  $A_1$



# Resilience boosting: details

Given:

Resynchroniser **B**

- $n = n_0 + n_1$
- $f = f_0 + f_1 + 1$

Pulser **A<sub>0</sub>**


Pulser **A<sub>1</sub>**

+

Wrapper **W**

Consensus **C**

- $n = n_0 + n_1$
- $f = f_0 + f_1$
- synchronous



# Resilience boosting: details

Given:

Resynchroniser **B**

- $n = n_0 + n_1$
- $f = f_0 + f_1 + 1$

Pulser **A<sub>0</sub>**

Pulser **A<sub>1</sub>**

+

Wrapper **W**

Consensus **C**

- $n = n_0 + n_1$
- $f = f_0 + f_1$
- synchronous

Result:

Pulser **A**





- $n = n_0 + n_1$  nodes
- $f = f_0 + f_1 + 1$  resilience

**B**

**A<sub>0</sub>**

**A<sub>1</sub>**


**C**



*arbitrary behaviour*

# Hardness of pulse synchronisation

**Synchronous BF consensus**



**SS + BF pulse synchronisation  
in the bounded delay model**

(almost)  
**as easy as**  
consensus

# Hardness of pulse synchronisation

**Synchronous BF consensus**

↑ ???




**SS + BF pulse synchronisation  
in the bounded delay model**

but is it  
as hard as  
consensus?



# Conclusions

## Upper bounds:




	stabilisation time	bandwidth
	<b>polylog <math>f</math></b>	<b>polylog <math>f</math></b>
	<b><math>O(\log f)</math></b>	<b>poly <math>f</math></b>
	<b><math>O(f)</math></b>	<b><math>O(\log f)</math></b>

## Lower bounds:

No lower bounds known!

# Conclusions

## Upper bounds:

	stabilisation time	bandwidth
	$\text{polylog } f$	$\text{polylog } f$
	$O(\log f)$	$\text{poly } f$
	$O(f)$	$O(\log f)$

## Lower bounds:

No lower bounds known!

**Thanks!**  
**arXiv:1705.06173**